

```
hadoop@pict-OptiPlex-5070:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
^Chadoop@pict-OptiPlex-5070:~$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [pict-OptiPlex-5070]
hadoop@pict-OptiPlex-5070:~$ jps
4352 SecondaryNameNode
4120 DataNode
4604 Jps
3965 NameNode
hadoop@pict-OptiPlex-5070:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hadoop@pict-OptiPlex-5070:~$ jps
4352 SecondaryNameNode
5330 Jps
4919 NodeManager
4120 DataNode
4749 ResourceManager
3965 NameNode
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -mkdir /m3
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -put /home/hadoop/test /m3/
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -ls /m3/
Found 1 items
-rw-r--r-- 1 hadoop supergroup      83 2023-05-05 14:46 /m3/test
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -cat /m3/test
hi
this is pict college
m3 batch is learning hadoop
hadoop is open source project
hadoop@pict-OptiPlex-5070:~$ yarn jar
/home/hadoop/hadoop-3.3.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.3.5.jar
wordcount /m3/test /m3_output
2023-05-05 14:51:11,208 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to
ResourceManager at /127.0.0.1:8032
2023-05-05 14:51:11,564 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for
path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1683277983499_0001
2023-05-05 14:51:11,755 INFO input.FileInputFormat: Total input files to process : 1
2023-05-05 14:51:12,014 INFO mapreduce.JobSubmitter: number of splits:1
2023-05-05 14:51:12,322 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_1683277983499_0001
2023-05-05 14:51:12,322 INFO mapreduce.JobSubmitter: Executing with tokens: []
2023-05-05 14:51:12,440 INFO conf.Configuration: resource-types.xml not found
2023-05-05 14:51:12,440 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2023-05-05 14:51:12,586 INFO impl.YarnClientImpl: Submitted application
application_1683277983499_0001
```

2023-05-05 14:51:12,643 INFO mapreduce.Job: The url to track the job: http://pict-OptiPlex-5070:8088/proxy/application_1683277983499_0001/
2023-05-05 14:51:12,643 INFO mapreduce.Job: Running job: job_1683277983499_0001
2023-05-05 14:51:18,711 INFO mapreduce.Job: Job job_1683277983499_0001 running in uber mode : false
2023-05-05 14:51:18,713 INFO mapreduce.Job: map 0% reduce 0%
2023-05-05 14:51:22,752 INFO mapreduce.Job: map 100% reduce 0%
2023-05-05 14:51:26,777 INFO mapreduce.Job: map 100% reduce 100%
2023-05-05 14:51:26,793 INFO mapreduce.Job: Job job_1683277983499_0001 completed successfully
2023-05-05 14:51:26,850 INFO mapreduce.Job: Counters: 54

File System Counters

FILE: Number of bytes read=147
FILE: Number of bytes written=552135
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=177
HDFS: Number of bytes written=93
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=1415
Total time spent by all reduces in occupied slots (ms)=1686
Total time spent by all map tasks (ms)=1415
Total time spent by all reduce tasks (ms)=1686
Total vcore-milliseconds taken by all map tasks=1415
Total vcore-milliseconds taken by all reduce tasks=1686
Total megabyte-milliseconds taken by all map tasks=1448960
Total megabyte-milliseconds taken by all reduce tasks=1726464

Map-Reduce Framework

Map input records=4
Map output records=15
Map output bytes=142
Map output materialized bytes=147
Input split bytes=94
Combine input records=15
Combine output records=12
Reduce input groups=12
Reduce shuffle bytes=147
Reduce input records=12
Reduce output records=12
Spilled Records=24
Shuffled Maps =1

Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=28
CPU time spent (ms)=750
Physical memory (bytes) snapshot=515792896
Virtual memory (bytes) snapshot=5587795968
Total committed heap usage (bytes)=392167424
Peak Map Physical memory (bytes)=276234240
Peak Map Virtual memory (bytes)=2791059456
Peak Reduce Physical memory (bytes)=239558656
Peak Reduce Virtual memory (bytes)=2796736512

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=83

File Output Format Counters

Bytes Written=93

```
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -ls /m3_outpt/
```

```
ls: `/m3_outpt/': No such file or directory
```

```
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -ls /m3_output/
```

```
Found 2 items
```

```
-rw-r--r--  1 hadoop supergroup      0 2023-05-05 14:51 /m3_output/_SUCCESS
```

```
-rw-r--r--  1 hadoop supergroup    93 2023-05-05 14:51 /m3_output/part-r-00000
```

```
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -ls /m3_output/part-r-00000
```

```
-rw-r--r--  1 hadoop supergroup    93 2023-05-05 14:51 /m3_output/part-r-00000
```

```
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -cat /m3_output/part-r-00000
```

```
batch  1
```

```
college 1
```

```
hadoop2
```

```
hi      1
```

```
is      3
```

```
learning 1
```

```
m3      1
```

```
open    1
```

```
pict    1
```

```
project 1
```

```
source  1
```

```
this    1
```

```
hadoop@pict-OptiPlex-5070:~$ stop-dfs.sh
```

```
Stopping namenodes on [localhost]
```

```
Stopping datanodes
```

```
Stopping secondary namenodes [pict-OptiPlex-5070]
```

```
hadoop@pict-OptiPlex-5070:~$ start-dfs.sh
```

```
Starting namenodes on [localhost]
```

```
Starting datanodes
Starting secondary namenodes [pict-OptiPlex-5070]
hadoop@pict-OptiPlex-5070:~$ jps
8337 DataNode
8595 SecondaryNameNode
8758 Jps
8119 NameNode
4919 NodeManager
4749 ResourceManager
hadoop@pict-OptiPlex-5070:~$ hdfs dfs -mkdir
```